

# a1

September 3, 2019

```
In [1]: import pandas as pd
import numpy as np
import altair as alt
import statsmodels.api as sm
import statsmodels.formula.api as smf
from math import log
from itertools import combinations
import copy
```

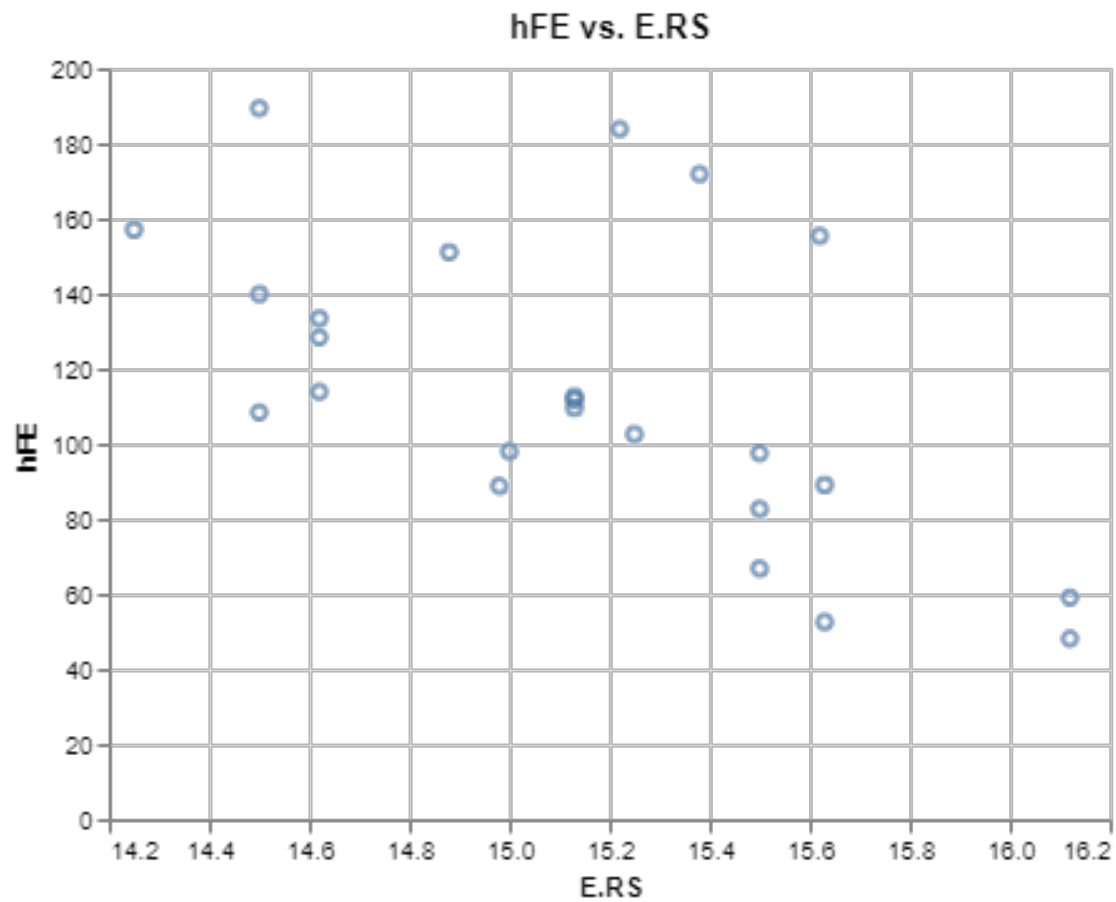
```
In [2]: semic = pd.read_csv('Data/semic.dat', sep='\s+')
        semic.columns=['ers', 'brs', 'ebrs', 'hfe']
        semic.head()
```

```
Out[2]:
```

	ers	brs	ebrs	hfe
0	14.62	226.0	7.000	128.40
1	15.63	220.0	3.375	52.62
2	14.62	217.4	6.375	113.90
3	15.00	220.0	6.000	98.01
4	14.50	226.5	7.625	139.90

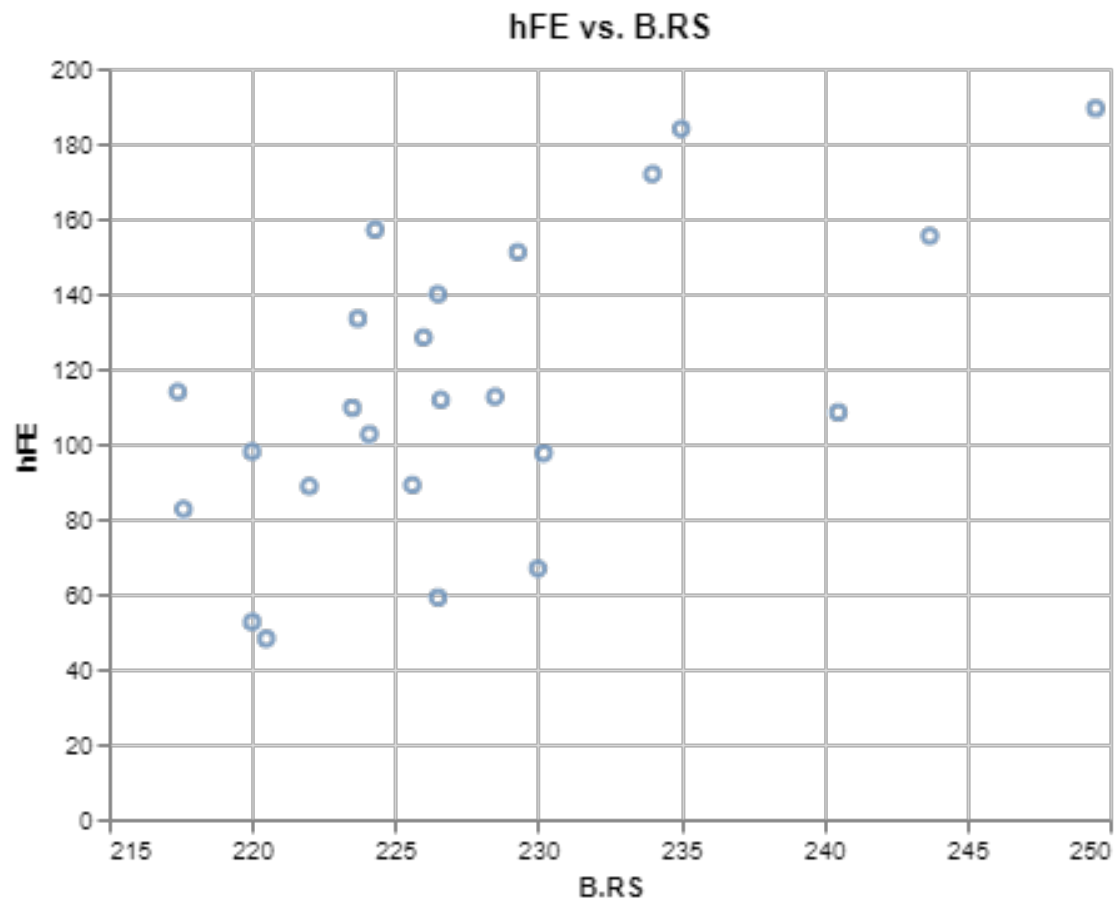
```
In [3]: alt.Chart(semic, title='hFE vs. E.RS').mark_point().encode(x=alt.X('ers', title='E.RS')
```

```
Out[3]:
```



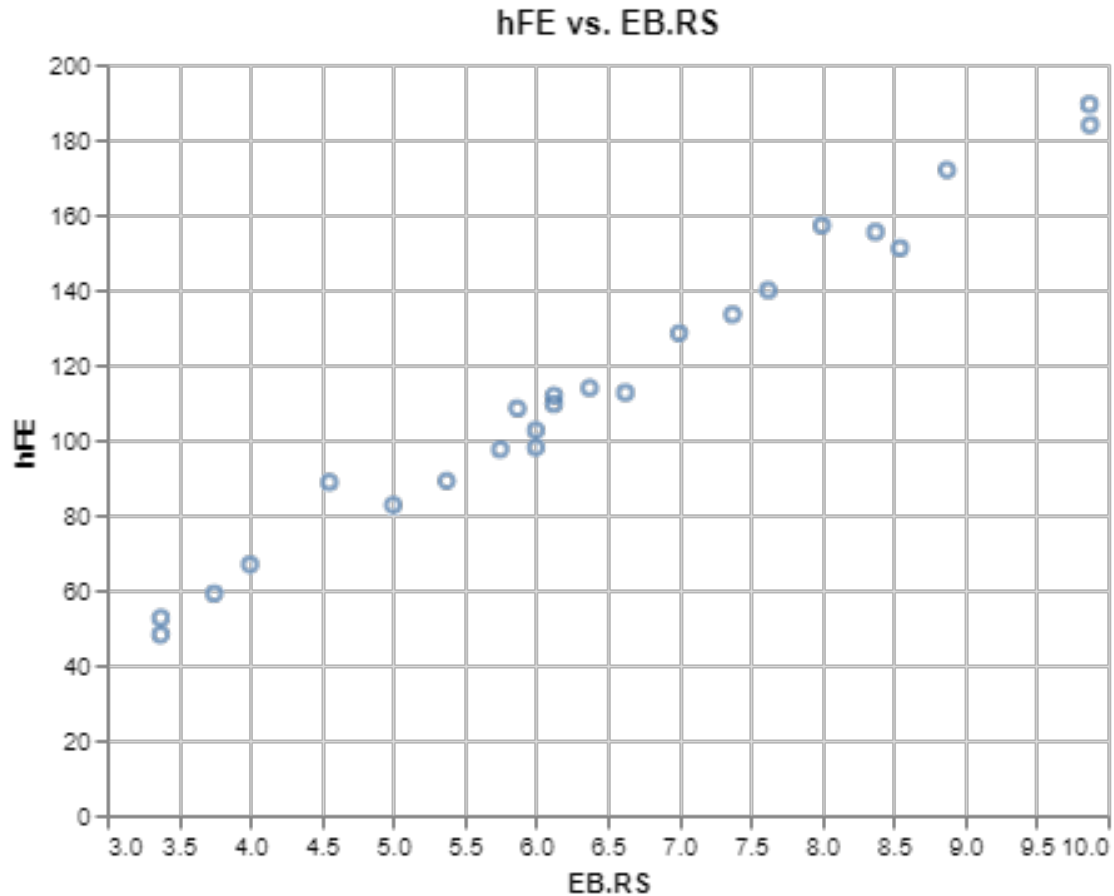
```
In [4]: alt.Chart(semic, title='hFE vs. B.RS').mark_point().encode(x=alt.X('brs', title='B.RS')
```

```
Out[4]:
```



```
In [5]: alt.Chart(semic, title='hFE vs. EB.RS').mark_point().encode(x=alt.X('ebrs', title='EB.RS'))
```

```
Out[5]:
```



```
In [6]: semicFit = smf.ols('hfe ~ ebrs + brs + ers', semic).fit()
print(semicFit.summary())
print(f'variance: {semicFit.scale}')
print(f'Coefficients: \n{semicFit.params}')
```

#### OLS Regression Results

```
=====
Dep. Variable:          hfe      R-squared:                0.988
Model:                  OLS      Adj. R-squared:           0.986
Method:                 Least Squares      F-statistic:          530.2
Date:                   Tue, 03 Sep 2019    Prob (F-statistic):      3.21e-19
Time:                   08:03:07    Log-Likelihood:         -69.287
No. Observations:       24      AIC:                   146.6
Df Residuals:           20      BIC:                   151.3
Df Model:                3
Covariance Type:        nonrobust
=====
```

```
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
```

Intercept	16.0097	44.187	0.362	0.721	-76.163	108.182
ebrs	19.4235	0.802	24.213	0.000	17.750	21.097
brs	0.2349	0.159	1.476	0.155	-0.097	0.567
ers	-5.2407	2.422	-2.164	0.043	-10.292	-0.189

---

Omnibus:	4.041	Durbin-Watson:	2.256
Prob(Omnibus):	0.133	Jarque-Bera (JB):	2.320
Skew:	0.693	Prob(JB):	0.314
Kurtosis:	3.630	Cond. No.	1.04e+04

---

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 1.04e+04. This might indicate that there are strong multicollinearity or other numerical problems.  
variance: 22.609981340360555

Coefficients:

Intercept 16.009655  
ebrs 19.423532  
brs 0.234928  
ers -5.240723  
dtype: float64

In [7]: `print(sm.stats.anova_lm(semicFit, type=2))`

	df	sum_sq	mean_sq	F	PR(>F)
ebrs	1.0	35840.470457	35840.470457	1585.161435	1.598031e-20
brs	1.0	19.231423	19.231423	0.850572	3.673831e-01
ers	1.0	105.899826	105.899826	4.683764	4.272500e-02
Residual	20.0	452.199627	22.609981	NaN	NaN

In [8]: `p = semicFit.params`  
`p[0] + 15.13*p[3] + 223.5*p[2] + 6.125*p[1]`

Out[8]: 108.19315918035909

In [9]: `p = semicFit.predict({'ebrs':6.125, 'brs':223.5, 'ers':15.13})`  
`p = semicFit.get_prediction({'ebrs':6.125, 'brs':223.5, 'ers':15.13})`

In [10]: `p.conf_int(alpha=.01)`

Out[10]: array([[105.01027994, 111.37603842]])

In [11]: `p.predicted_mean`

Out[11]: array([108.19315918])

In [12]: `p.conf_int(obs=True, alpha=0.01)`

```
Out[12]: array([[ 94.29423185, 122.09208651]])
```

```
In [13]: semicFit.conf_int(.01)
```

```
Out[13]:
```

	0	1
Intercept	-109.717212	141.736523
ebrs	17.141004	21.706059
brs	-0.217842	0.687699
ers	-12.130860	1.649413

```
In [14]: semicFit.cov_params()
```

```
Out[14]:
```

	Intercept	ebrs	brs	ers
Intercept	1952.486049	-3.435156	-3.671093	-72.237543
ebrs	-3.435156	0.643523	-0.079168	1.145233
brs	-3.671093	-0.079168	0.025321	-0.104852
ers	-72.237543	1.145233	-0.104852	5.863911

```
In [15]: semic.mean()
```

```
Out[15]:
```

ers	15.138750
brs	227.708333
ebrs	6.410000
hfe	114.671667

dtype: float64

```
In [16]: semic.head()
```

```
Out[16]:
```

	ers	brs	ebrs	hfe
0	14.62	226.0	7.000	128.40
1	15.63	220.0	3.375	52.62
2	14.62	217.4	6.375	113.90
3	15.00	220.0	6.000	98.01
4	14.50	226.5	7.625	139.90

```
In [17]: semicCntr = semic/semic.mean()  
semicCntr.head()
```

```
Out[17]:
```

	ers	brs	ebrs	hfe
0	0.965734	0.992498	1.092044	1.119719
1	1.032450	0.966148	0.526521	0.458875
2	0.965734	0.954730	0.994540	0.993271
3	0.990835	0.966148	0.936037	0.854701
4	0.957807	0.994694	1.189548	1.220005

```
In [18]: X = np.array(semicCntr[['ers','brs','ebrs']])  
S = np.matmul(X.transpose(), X)  
S
```

```
Out[18]: array([[24.02636101, 23.99594666, 23.87148981],
               [23.99594666, 24.02860805, 24.14206085],
               [23.87148981, 24.14206085, 26.0160813 ]])
```

```
In [19]: S.diagonal()*0.5
```

```
Out[19]: array([4.90166921, 4.90189841, 5.10059617])
```

```
In [20]: S
```

```
Out[20]: array([[24.02636101, 23.99594666, 23.87148981],
               [23.99594666, 24.02860805, 24.14206085],
               [23.87148981, 24.14206085, 26.0160813 ]])
```

```
In [21]: X[0,1]
```

```
Out[21]: 0.9924977127172918
```

```
In [22]: X[23]
```

```
Out[22]: array([1.03178928, 1.07022873, 1.30655226])
```

```
In [23]: r = np.zeros([3,3])
         for i in range(3):
             for j in range(3):
                 r[i,j] = S[i,j]/(np.sqrt(S[i,i]*S[j,j]))
         print(r)
```

```
[[1.          0.99868743 0.95480478]
 [0.99868743 1.          0.96558185]
 [0.95480478 0.96558185 1.          ]]
```

## 0.1 Question 2

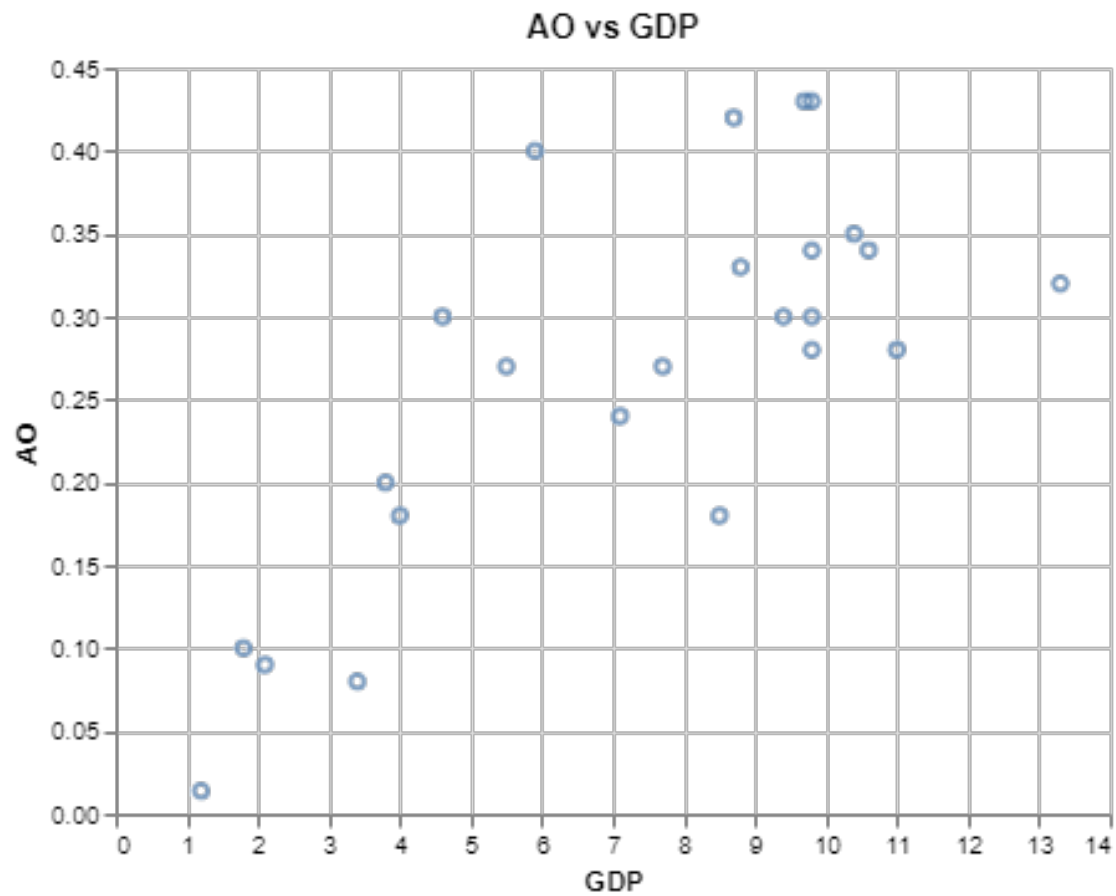
```
In [24]: cars = pd.read_csv('Data/car.dat', sep='\s+')
         cars.head()
```

```
Out[24]:
```

	Country	AO	POP	DEN	GDP	PR	CON	TR
0	Austria	0.27	7.5	89	7.7	49	1.11	2.6
1	Belgium	0.30	9.8	323	9.8	59	1.04	1.6
2	Canada	0.42	23.5	2	8.7	17	2.82	0.1
3	Denmark	0.28	5.1	119	11.0	56	1.21	1.9
4	Finland	0.24	4.8	16	7.1	49	1.22	2.2

```
In [25]: alt.Chart(cars, title='AO vs GDP').mark_point().encode(x='GDP', y=alt.Y('AO', scale=a
```

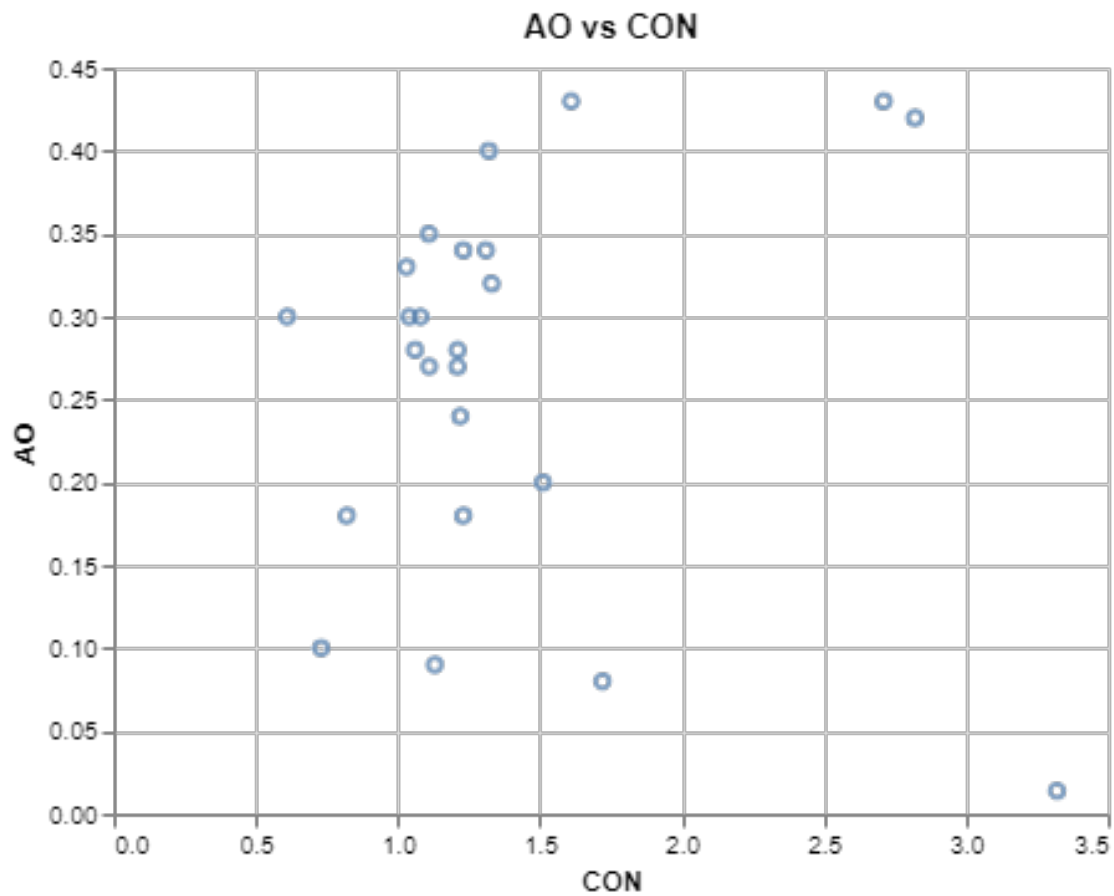
```
Out[25]:
```



```
In [26]: alt.Chart(cars, title='AO vs CON').mark_point().encode(x='CON', y=alt.Y('AO', scale=alt.Scale(zero=False)))
```

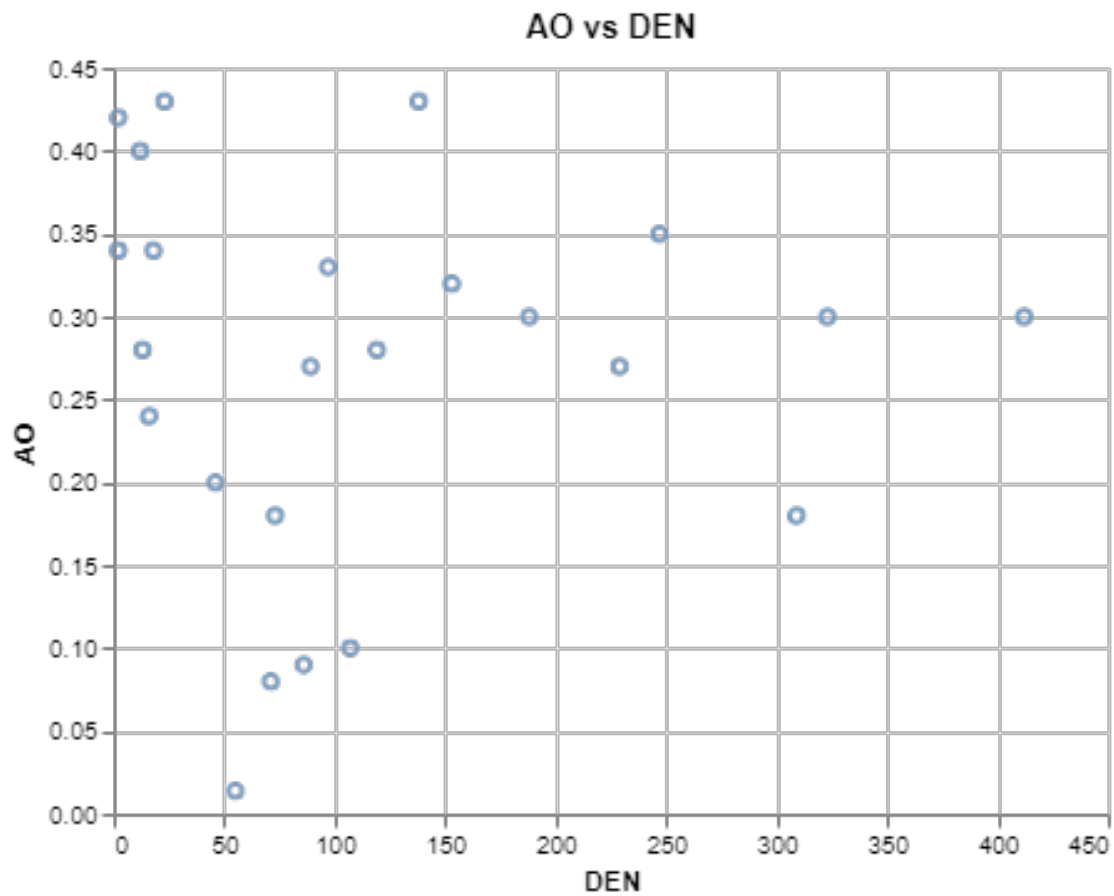
Out[26]:





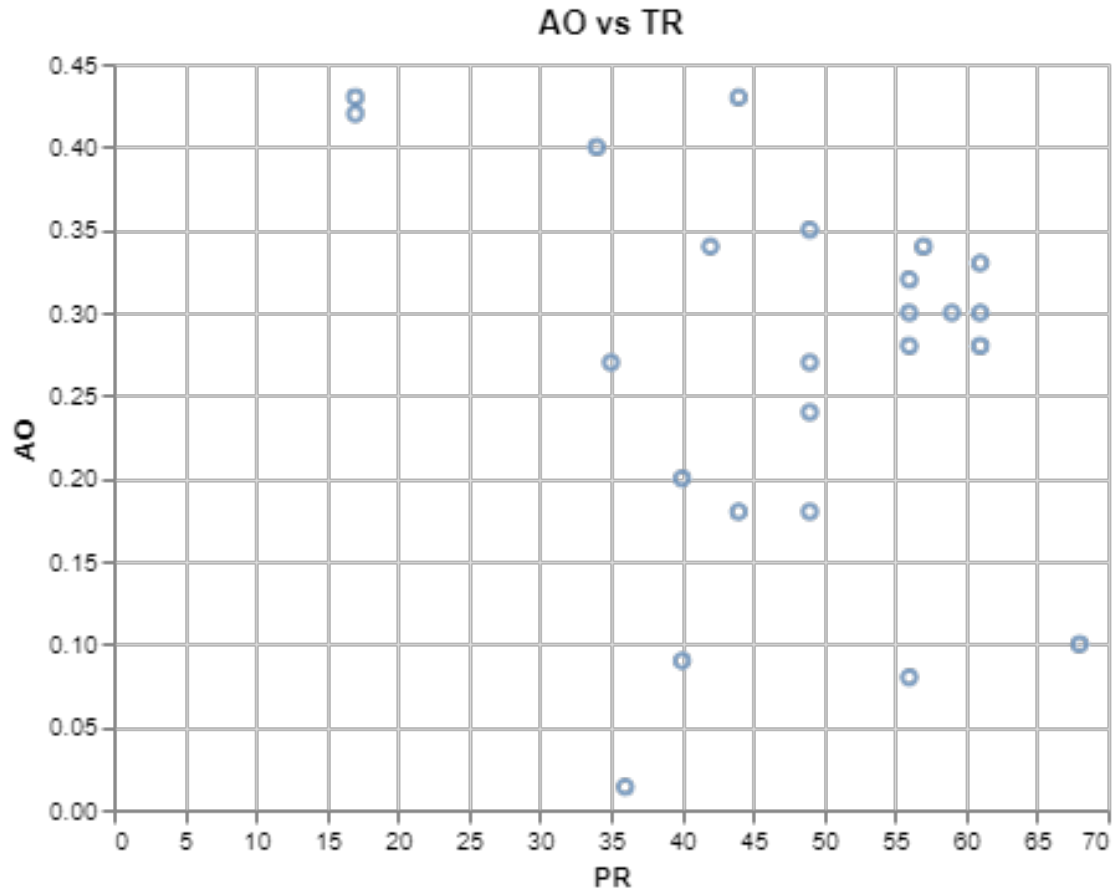
```
In [27]: alt.Chart(cars, title='AO vs DEN').mark_point().encode(x='DEN', y=alt.Y('AO', scale=a
```

```
Out[27]:
```



```
In [28]: alt.Chart(cars, title='AO vs TR').mark_point().encode(x='PR', y=alt.Y('AO', scale=alt
```

```
Out[28]:
```



```
In [29]: #code for forward_selected from https://planspace.org/20150423-forward\_selection\_with
def forward_selected(data, response):
    """Linear model designed by forward selection.

    Parameters:
    -----
    data : pandas DataFrame with all possible predictors and response

    response: string, name of response column in data

    Returns:
    -----
    model: an "optimal" fitted statsmodels linear model
           with an intercept
           selected by forward selection
           evaluated by adjusted R-squared
    """
    remaining = set(data.columns)
    remaining.remove(response)
```

```

selected = []
current_score, best_new_score = 0.0, 0.0
while remaining and current_score == best_new_score:
    scores_with_candidates = []
    for candidate in remaining:
        formula = "{} ~ {} + 1".format(response,
                                         ' + '.join(selected + [candidate]))
        score = smf.ols(formula, data).fit().rsquared_adj
        scores_with_candidates.append((score, candidate))
    scores_with_candidates.sort()
    best_new_score, best_candidate = scores_with_candidates.pop()
    if current_score < best_new_score:
        remaining.remove(best_candidate)
        selected.append(best_candidate)
        current_score = best_new_score
formula = "{} ~ {} + 1".format(response,
                                ' + '.join(selected))
model = smf.ols(formula, data).fit()
return model

```

```

In [47]: def backward_selected(data, response):
    cols = set(data.columns)
    cols.remove(response)
    cols = list(cols)
    currentScore, bestNewScore = 0.0, 0.0
    loop = True

    while loop and currentScore == bestNewScore:
        scoresAndCandidates = []
        rCols = copy.deepcopy(cols)
        for col in rCols:
            rCols.remove(col)
            formula = "{} ~ {}".format(response, ' + '.join(rCols))
            score = smf.ols(formula, data).fit().rsquared_adj
            scoresAndCandidates.append((score, col))
        scoresAndCandidates.sort()
        bestNewScore, bestCandidate = scoresAndCandidates.pop()
        if currentScore < bestNewScore:
            cols.remove(bestCandidate)
            currentScore = bestNewScore
        else:
            loop = False
    formula = f"{response} ~ {' + '.join(cols)}"
    model = smf.ols(formula, data).fit()
    return model

In [48]: m = backward_selected(cars[['AO', 'POP', 'DEN', 'GDP', 'PR', 'CON', 'TR']], 'AO')
print(m.summary())
print(f'Variance: {m.scale}')

```

### OLS Regression Results

```
=====
Dep. Variable:          AO      R-squared:            0.843
Model:                  OLS      Adj. R-squared:        0.799
Method:                 Least Squares      F-statistic:      19.28
Date:                  Tue, 03 Sep 2019      Prob (F-statistic):  1.16e-06
Time:                  08:07:55      Log-Likelihood:      40.693
No. Observations:      24      AIC:                -69.39
Df Residuals:          18      BIC:                -62.32
Df Model:               5
Covariance Type:       nonrobust
=====
```

```
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept      0.4834      0.089       5.412      0.000      0.296      0.671
CON            -0.1007      0.026      -3.881      0.001     -0.155     -0.046
DEN           -3.403e-05      0.000     -0.323      0.750     -0.000      0.000
PR             -0.0044      0.001     -3.773      0.001     -0.007     -0.002
TR            -0.0612      0.017     -3.580      0.002     -0.097     -0.025
GDP            0.0308      0.003      8.932      0.000      0.024      0.038
=====
```

```
=====
Omnibus:          3.592      Durbin-Watson:          2.322
Prob(Omnibus):    0.166      Jarque-Bera (JB):          2.081
Skew:             0.686      Prob(JB):                 0.353
Kurtosis:         3.446      Cond. No.                  1.47e+03
=====
```

### Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
 [2] The condition number is large, 1.47e+03. This might indicate that there are strong multicollinearity or other numerical problems.  
 Variance: 0.002628573282679365

```
In [33]: m = forward_selected(cars[['AO', 'POP', 'DEN', 'GDP', 'PR', 'CON', 'TR']], 'AO')
         print(m.summary())
         print(f'Variance: {m.scale}')
```

### OLS Regression Results

```
=====
Dep. Variable:          AO      R-squared:            0.842
Model:                  OLS      Adj. R-squared:        0.808
Method:                 Least Squares      F-statistic:      25.26
Date:                  Tue, 03 Sep 2019      Prob (F-statistic):  2.23e-07
Time:                  08:05:19      Log-Likelihood:      40.624
No. Observations:      24      AIC:                -71.25
Df Residuals:          19      BIC:                -65.36
Df Model:               4
=====
```

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.4831	0.087	5.541	0.000	0.301	0.666
GDP	0.0307	0.003	9.163	0.000	0.024	0.038
TR	-0.0623	0.016	-3.804	0.001	-0.097	-0.028
CON	-0.1001	0.025	-3.963	0.001	-0.153	-0.047
PR	-0.0045	0.001	-3.930	0.001	-0.007	-0.002
Omnibus:		3.521	Durbin-Watson:			2.434
Prob(Omnibus):		0.172	Jarque-Bera (JB):			2.151
Skew:		0.717	Prob(JB):			0.341
Kurtosis:		3.310	Cond. No.			437.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
Variance: 0.0025046815666849172

```
In [34]: def calcPress(X, Y, Ypred):
          X = np.matrix(X)
          h = np.matmul(X, np.matmul(np.linalg.inv(np.matmul(X.transpose(), X)), X.transpose())
          press = 0
          #for i in range(lenh(Y)):
          press = press + (((Y - Ypred)**2 / (1-np.diagonal(h))))**2).sum()
          return press

In [35]: def calcYpred(model, inputs):
          y = model.predict(inputs)
          return y

In [36]: np.diagonal([[1,2],[3,4]])

Out[36]: array([1, 4])

In [37]: %%time
          variables = ['GDP', 'PR', 'CON', 'TR']
          #variables = variables.drop('y')
          bestFormula = ''
          bestAIC = 9999999
          maxAIC = 0

          fullModelMSE = smf.ols('AO ~ POP + DEN + GDP + PR + CON + TR', cars).fit().mse_resid

          for i in range(1, len(variables)+1):
              combs = combinations(variables, i)
```

```

for c in combs:
    formula = 'AO ~'
    for v in c:
        formula = formula + f' + {v}'
    formula = formula.replace('~ + ', '~ ')
    fit = smf.ols(formula, cars).fit()
    fit.predict()
    if fit.aic < bestAIC:
        bestAIC = fit.aic
        bestFormula = formula
    if fit.aic > maxAIC:
        maxAIC = fit.aic
    ypred = calcYpred(fit, cars[[x for x in c]])
    print(f'{formula} : ')
    press = str(calcPress(cars[[x for x in c]], cars.AO, ypred))[:6]
    cm = ((fit.mse_resid * fit.df_resid) / fullModelMSE) - len(cars) + 2*len(fit.)

    print(f'AIC: {str(fit.aic)[:8]} - PRESS: {press} - RSquared: {str(fit.rsquared)}')

```

```

AO ~ GDP :
AIC: -52.8908 - PRESS: 0.0019 - RSquared: 0.5633 - m: 2 - Cm: 29.026
AO ~ PR :
AIC: -34.7088 - PRESS: 0.0105 - RSquared: 0.0685 - m: 2 - Cm: 84.578
AO ~ CON :
AIC: -33.0052 - PRESS: 0.0127 - RSquared: 6.7618 - m: 2 - Cm: 92.271
AO ~ TR :
AIC: -33.1466 - PRESS: 0.0105 - RSquared: 0.0058 - m: 2 - Cm: 91.612
AO ~ GDP + PR :
AIC: -57.2986 - PRESS: 0.0015 - RSquared: 0.6656 - m: 3 - Cm: 19.538
AO ~ GDP + CON :
AIC: -51.1158 - PRESS: 0.0026 - RSquared: 0.5674 - m: 3 - Cm: 30.568
AO ~ GDP + TR :
AIC: -58.3450 - PRESS: 0.0012 - RSquared: 0.6799 - m: 3 - Cm: 17.936
AO ~ PR + CON :
AIC: -34.5919 - PRESS: 0.0079 - RSquared: 0.1388 - m: 3 - Cm: 78.686
AO ~ PR + TR :
AIC: -32.7925 - PRESS: 0.0115 - RSquared: 0.0717 - m: 3 - Cm: 86.214
AO ~ CON + TR :
AIC: -31.2181 - PRESS: 0.0158 - RSquared: 0.0088 - m: 3 - Cm: 93.280
AO ~ GDP + PR + CON :
AIC: -59.6559 - PRESS: 0.0009 - RSquared: 0.7211 - m: 4 - Cm: 15.305
AO ~ GDP + PR + TR :
AIC: -58.7871 - PRESS: 0.0015 - RSquared: 0.7108 - m: 4 - Cm: 16.459
AO ~ GDP + CON + TR :
AIC: -58.9709 - PRESS: 0.0008 - RSquared: 0.7130 - m: 4 - Cm: 16.212
AO ~ PR + CON + TR :
AIC: -32.6888 - PRESS: 0.0086 - RSquared: 0.1422 - m: 4 - Cm: 80.297

```

```
AO ~ GDP + PR + CON + TR :
AIC: -71.2479 - PRESS: 0.0004 - RSquared: 0.8417 - m: 5 - Cm: 3.7693
Wall time: 589 ms
```

```
In [38]: print(bestFormula)
         print(bestAIC)
         print(maxAIC)
```

```
AO ~ GDP + PR + CON + TR
-71.2479549532554
0
```

## 0.2 Question 3

```
In [39]: jelly = pd.read_csv('Data/jelly.dat', sep='\s+')
         jelly['Site Name'] = jelly.Site.map(lambda x: 'Salamander Bay' if x==2 else 'Dangar Island')
         jelly.head()
```

```
Out[39]:
```

	Breadth	Length	Site	Site Name
0	6.5	8.5	1	Dangar Island
1	6.0	9.0	1	Dangar Island
2	6.5	9.0	1	Dangar Island
3	7.0	9.0	1	Dangar Island
4	8.0	9.5	1	Dangar Island

```
In [40]: j1Fit = smf.ols('Breadth ~ Length', jelly[jelly.Site==1]).fit()
         print(j1Fit.summary())
```

```

OLS Regression Results
=====
Dep. Variable:          Breadth      R-squared:                0.915
Model:                  OLS          Adj. R-squared:           0.910
Method:                 Least Squares  F-statistic:              214.6
Date:                  Tue, 03 Sep 2019  Prob (F-statistic):       3.72e-12
Time:                  08:05:56       Log-Likelihood:           -29.794
No. Observations:      22           AIC:                       63.59
Df Residuals:          20           BIC:                       65.77
Df Model:              1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-2.9374	0.920	-3.194	0.005	-4.856	-1.019
Length	1.0534	0.072	14.650	0.000	0.903	1.203

```
=====
Omnibus:                0.410    Durbin-Watson:              2.254
Prob(Omnibus):          0.815    Jarque-Bera (JB):          0.043
```



```
Skew: -0.108 Prob(JB): 0.979
Kurtosis: 3.005 Cond. No. 56.5
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [41]: j2Fit = smf.ols('Breadth ~ Length', jelly[jelly.Site==2]).fit()
print(j2Fit.summary())
```

```

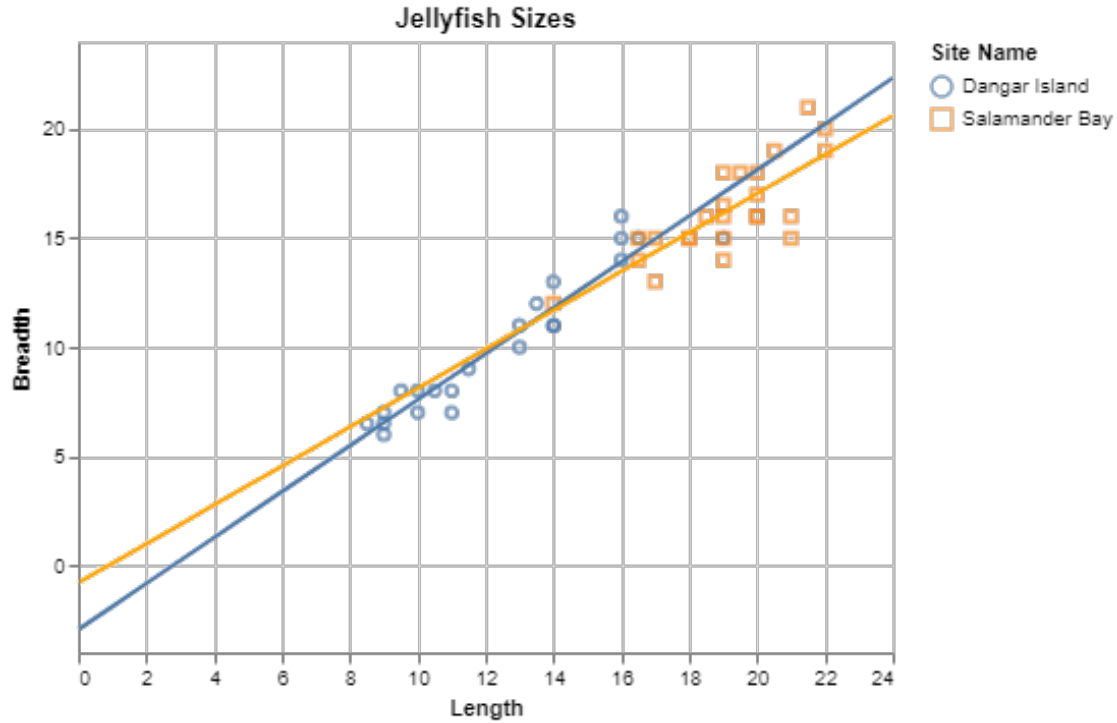
                        OLS Regression Results
=====
Dep. Variable:          Breadth    R-squared:                0.625
Model:                  OLS        Adj. R-squared:            0.608
Method:                 Least Squares    F-statistic:            36.68
Date:                  Tue, 03 Sep 2019    Prob (F-statistic):      4.28e-06
Time:                  08:05:56    Log-Likelihood:          -40.551
No. Observations:      24    AIC:                        85.10
Df Residuals:          22    BIC:                        87.46
Df Model:               1
Covariance Type:       nonrobust
=====
                        coef    std err          t      P>|t|      [0.025    0.975]
-----
Intercept             -0.8003      2.826     -0.283      0.780     -6.661     5.060
Length                 0.8924      0.147      6.056      0.000      0.587     1.198
=====
Omnibus:               0.436    Durbin-Watson:           1.849
Prob(Omnibus):         0.804    Jarque-Bera (JB):         0.405
Skew:                  -0.274    Prob(JB):                 0.817
Kurtosis:              2.676    Cond. No.:                194.
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [42]: alt.Chart(jelly, title='Jellyfish Sizes').mark_point().encode(x='Length', y='Breadth')
alt.Chart(pd.DataFrame({'Breadth': [j1Fit.params[0] + x*j1Fit.params[1] for x in range(25)]},
                        {'Length': [x for x in range(25)]})).mark_line().encode(x='Length', y='Breadth')
alt.Chart(pd.DataFrame({'Breadth': [j2Fit.params[0] + x*j2Fit.params[1] for x in range(25)]},
                        {'Length': [x for x in range(25)]})).mark_line(color='orange').encode(x='Length', y='Breadth')
```

Out[42]:



### 0.3 Question 4

```
In [43]: cloud = pd.read_csv('Data\cloud.dat', sep='\s+')
         #cloud.A = cloud.A + 1
         for col in ['S', 'C', 'P', 'E']:
             cloud[f'A{col}'] = cloud.A * cloud[col]
         cloud.head()
```

```
Out[43]:
```

	A	T	S	C	P	E	y	AS	AC	AP	AE
0	0	0	1.75	13.4	0.274	2	2.61	0.00	0.0	0.000	0
1	1	3	4.10	3.9	0.198	2	1.81	4.10	3.9	0.198	2
2	0	4	2.35	5.3	0.526	1	1.78	0.00	0.0	0.000	0
3	1	6	4.25	7.1	0.250	1	0.83	4.25	7.1	0.250	1
4	0	9	1.60	6.9	0.018	2	1.28	0.00	0.0	0.000	0

```
In [44]: cFit = smf.ols('y ~ A + T + S + C + P + E + AS + AC + AP + AE', cloud).fit()
         print(cFit.summary())
         print(cFit.scale)
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y      R-squared:                0.879
Model:                            OLS      Adj. R-squared:         0.744
Method:                           Least Squares      F-statistic:           6.518

```

Date: Tue, 03 Sep 2019 Prob (F-statistic): 0.00473  
Time: 08:05:57 Log-Likelihood: -0.43477  
No. Observations: 20 AIC: 22.87  
Df Residuals: 9 BIC: 33.82  
Df Model: 10  
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-0.3469	0.971	-0.357	0.729	-2.544	1.850
A	3.4404	1.213	2.835	0.020	0.695	6.185
T	-0.0147	0.006	-2.435	0.038	-0.028	-0.001
S	0.1267	0.226	0.560	0.589	-0.385	0.639
C	0.0746	0.041	1.825	0.101	-0.018	0.167
P	1.1913	0.677	1.759	0.112	-0.341	2.723
E	0.6372	0.367	1.735	0.117	-0.194	1.468
AS	-0.8519	0.277	-3.077	0.013	-1.478	-0.226
AC	-0.0588	0.078	-0.751	0.472	-0.236	0.118
AP	0.5972	3.313	0.180	0.861	-6.898	8.092
AE	0.0675	0.483	0.140	0.892	-1.025	1.160
Omnibus:	0.629	Durbin-Watson:	2.691			
Prob(Omnibus):	0.730	Jarque-Bera (JB):	0.586			
Skew:	0.357	Prob(JB):	0.746			
Kurtosis:	2.562	Cond. No.	1.70e+03			

#### Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 1.7e+03. This might indicate that there are strong multicollinearity or other numerical problems.  
0.13589230416186726

In [45]: s = '='

```
for p in ['A' , 'T' , 'S' , 'C' , 'P' , 'E' , 'AS' , 'AC' , 'AP' , 'AE']:
    symbol = '+' if cFit.params[p] >=0 else '-'
    s = s + f'{symbol}{str(cFit.params[p])[:6]}*{p} '

s
```

Out[45]: '= +3.4404\*A -0.014\*T +0.1267\*S +0.0745\*C +1.1913\*P +0.6372\*E -0.851\*AS -0.058\*AC +0.0675\*AE'

```
In [46]: cloud['residuals'] = cloud.y - cyPred
cloud['normResiduals'] = cFit.get_influence().resid_studentized_internal
cloud['yFit'] = cyPred
```

NameError

Traceback (most recent call last)

```
<ipython-input-46-b15db7622b86> in <module>()
----> 1 cloud['residuals'] = cloud.y - cyPred
      2 cloud['normResiduals'] = cFit.get_influence().resid_studentized_internal
      3 cloud['yFit'] = cyPred
```

NameError: name 'cyPred' is not defined

```
In [ ]: cFit.fvalue
```

```
In [ ]: cFit.get_influence().resid_studentized_internal
```

```
In [ ]: sm.qqplot(cFit.get_influence().resid_studentized_internal, line='45')
```

```
In [ ]: from statsmodels.graphics.gofplots import ProbPlot as pPlot
        from matplotlib import pyplot as plt
```

```
In [ ]: %matplotlib inline
        fig = pPlot(cyPred-cloud.y)
        plt.show()
```

```
In [ ]: cloud.head()
```

```
In [ ]: alt.Chart(cloud).mark_point().encode(x='yFit', y='normResiduals')
```

```
In [ ]: alt.Chart(cloud).mark_point().encode(x='A', y='normResiduals')
```

```
In [ ]: alt.Chart(cloud).mark_point().encode(x='AS', y='normResiduals')
```